

EMNLP

Meaning Representation and Semantic Analysis

Yansong Feng
fengyansong@pku.edu.cn

Institute of Computer Science and Technology
Peking University

June 12, 2018

Recap: Semantics

- Distributional semantics
- Shallow semantics

Recap: Semantics

- Distributional semantics
- Shallow semantics
- Less clear how to deal with compositionality
- Still haven't discussed how to do inference

- Question
Did Poland reduce its carbon emissions since 1989?
- Resources available
Due to the collapse of the industrial sector after the end of communism in 1989, all countries in Central Europe saw a fall in carbon emissions. Poland is a country in Central Europe.

- Question
Did Poland reduce its carbon emissions since 1989?
- Resources available
Due to the collapse of the industrial sector after the end of communism in 1989, all countries in Central Europe saw a fall in carbon emissions. Poland is a country in Central Europe.
- What is hard?

- Question
Did Poland reduce its carbon emissions since 1989?
- Resources available
Due to the collapse of the industrial sector after the end of communism in 1989, all countries in Central Europe saw a fall in carbon emissions. Poland is a country in Central Europe.
- What is hard?
 - need to do inference
 - a problem for sentential, not lexical, semantics

- **Vector Space Model** is one kind of meaning representation (MR)
- For **compositionality** and **inference**, we need meaning representations that are **symbolic** and **structured**.
- More difficult: how about the MRs of sentences (using syntax to help).

Meaning Representations

- Vector Space Model is one kind of meaning representation (MR)
- For **compositionality** and **inference**, we need meaning representations that are **symbolic** and **structured**.
- More difficult: how about the MRs of sentences (using syntax to help).

Firstly, define the semantics we aim at, i.e., a **meaning representation language (MRL)**.

The symbols in our meaning representations correspond to objects, properties, and relations in the world.

- The world may be the real world, or (usually) a formalized and well-specified world: a model or knowledge base of known facts.
 - a tiny world model containing 3 entities, and an exhaustive table of ‘who loves whom’ relations.
 - GeoQuery dataset, containing 800 facts about US geography.
 - Freebase, “A community-curated database of well-known people, places, and things” with over 2.6 billion facts.

What do we expect?

- **Compositional**: The meaning of a complex expression is a function of the meaning of its parts and of the rules by which they are combined.
- **Verifiable**: Can use the MR of a sentence to determine whether the sentence is true with respect to some given model of the world.
- **Unambiguous**: an MR should have exactly one interpretation. So, an ambiguous sentence should have a different MR for each sense.
- **Canonical form**: Sentences with the same (literal) meaning should have the same MR.
- **Inference**: Be able to verify sentences not only directly, but also by drawing conclusions based on the input MR and facts in the KB.
- **Expressivity**: Allow us to handle a wide range of meanings and express appropriate relationships between the words in a sentence.

What do we expect?

- **Compositional**: The meaning of a complex expression is a function of the meaning of its parts and of the rules by which they are combined.
- **Verifiable**: Can use the MR of a sentence to determine whether the sentence is true with respect to some given model of the world.
- **Unambiguous**: an MR should have exactly one interpretation. So, an ambiguous sentence should have a different MR for each sense.
 - each interpretation of *time flies like an arrow* should have a distinct MR
- **Canonical form**: Sentences with the same (literal) meaning should have the same MR.
- **Inference**: Be able to verify sentences not only directly, but also by drawing conclusions based on the input MR and facts in the KB.
- **Expressivity**: Allow us to handle a wide range of meanings and express appropriate relationships between the words in a sentence.

What do we expect?

- **Compositional**: The meaning of a complex expression is a function of the meaning of its parts and of the rules by which they are combined.
- **Verifiable**: Can use the MR of a sentence to determine whether the sentence is true with respect to some given model of the world.
- **Unambiguous**: an MR should have exactly one interpretation. So, an ambiguous sentence should have a different MR for each sense.
- **Canonical form**: Sentences with the same (literal) meaning should have the same MR.
 - Tanjore serves vegetarian food and Vegetarian dishes are served by Tanjore should have the same canonical form
- **Inference**: Be able to verify sentences not only directly, but also by drawing conclusions based on the input MR and facts in the KB.
- **Expressivity**: Allow us to handle a wide range of meanings and express appropriate relationships between the words in a sentence.

What do we expect?

- **Compositional**: The meaning of a complex expression is a function of the meaning of its parts and of the rules by which they are combined.
- **Verifiable**: Can use the MR of a sentence to determine whether the sentence is true with respect to some given model of the world.
- **Unambiguous**: an MR should have exactly one interpretation. So, an ambiguous sentence should have a different MR for each sense.
- **Canonical form**: Sentences with the same (literal) meaning should have the same MR.
- **Inference**: Be able to verify sentences not only directly, but also by drawing conclusions based on the input MR and facts in the KB.
 - query: Did Poland reduce its carbon emissions?
 - facts: Carbon emissions have fallen for all countries in Central Europe. and Poland is a country in Central Europe.
 - answer: yes
- **Expressivity**: Allow us to handle a wide range of meanings and express appropriate relationships between the words in a sentence.

- Predicate Logic
- A simpler MRL that covers a lot of what we want.
- $tall(Kim) \vee tall(Pierre)$
- $likes(Sam, ownerOf(Tanjore))$
- $\exists x.cat(x) \wedge owns(Marie, x)$
- $\exists x.movie(x) \wedge \forall y.person(y) \Rightarrow loves(y, x)$

- Expressions are constructed from terms:
 - constant and variable symbols that represent entities
 - function symbols that allow us to indirectly specify entities
 - predicate symbols that represent properties of entities and relations between entities

- Expressions are constructed from terms:
 - constant and variable symbols that represent entities
 - function symbols that allow us to indirectly specify entities
 - predicate symbols that represent properties of entities and relations between entities
- Terms can be combined into predicate-argument structures, which in turn are combined into complex expressions using:
 - Logical connectives:
 - Quantifiers: \forall (universal quantifier, i.e., “for all”), \exists (existential quantifier, i.e. “exists”)

Constants

- Each constant symbol denotes exactly one entity
- Not all entities have a constant that denotes them
- Several constant symbols may denote the same entity:

Predicates

- Predicates with one argument represent properties of entities:
`nation(Scotland)`, `organization(EU)`, `tall(John)`
- Predicates with multiple arguments represent relations between entities:
`member-of(UK, EU)`, `likes(John, Marie)`, `introduced(John, Marie, Sue)`
- We write “/N” to indicate that a predicate has arity N (takes N arguments)
`member-of/2`, `nation/1`, `tall/1`, `introduced/3`

The Semantics of Predicates and Functions

Predicates:

- A predicate of arity N denotes the set of N -tuples that satisfy it.
likes/2 is the set of (x, y) pairs for which likes(x, y) is true
- If all arguments are instantiated, then the predicate-argument structure has a **truth value** (determined by comparing it to the set of facts in the world).
nation(UK) is true, locatedIn(China, Europe) is not true

Functions:

- look like unary predicates
- used to specify (denote) unique entities **indirectly**
- *president(EU), father(John)*

Logical Connectives, Variables, Quantifiers

Logical Connectives

- $\neg, \vee, \wedge, \Rightarrow$

Variables

- range over entities
- An expression consisting only of a predicate with a variable among its arguments is interpreted as a set:
likes(x, Kim) is the set of entities that like Kim
- A predicate with a variable among its arguments only has a truth value if it is bound by a quantifier.
 $\forall x.likes(x, Kim)$ has an interpretation as either true or false.

Quantifiers

- Universal Quantifier \forall :
Cats are mammals has MR $\forall x.cat(x) \Rightarrow mammal(x)$
- Existential Quantifier \exists :
Marie owns a cat has MR $\exists x.cat(x) \wedge owns(Marie, x)$

Advantages

- verifiable, unambiguous, canonical
- Predicate-argument structure
- Determiners and coordination → logical connectives and quantifiers

- verifiable, unambiguous, canonical
- Predicate-argument structure
- Determiners and coordination → logical connectives and quantifiers
- Compositionality??

Lambda (λ) Expressions

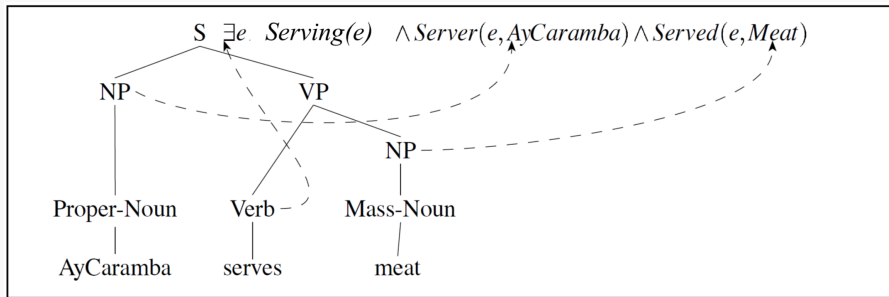
- Extension to FOL, allows us to work with ‘partially constructed’ formulae
- A λ -expression consists of:
 - the Greek letter λ , followed by a **formal parameter**
 - a FOL expression that may involve that variable
 $\lambda x.nation(x)$
- λ -reduction:
 - *$\lambda x.sleep(x)(Marie)$ becomes $sleep(Marie)$*
 - *$\lambda y.\lambda x.love(x, y)(pizza)$ becomes $\lambda x.love(x, pizza)$*
 - *$\lambda x.love(x, pizza)(Marie)$ becomes $love(Marie, pizza)$*

- **Reifying** events: introduce variables for events, which we can quantify over
- John gave Mary a book
- $\exists e, z. Giving(e) \wedge Giver(e, John) \wedge Givee(e, Mary) \wedge Given(e, z) \wedge Book(z)$
- John gave Mary a book on Tuesday
- $\exists e, z. Giving(e) \wedge Giver(e, John) \wedge Givee(e, Mary) \wedge Given(e, z) \wedge Book(z) \wedge Time(e, Tuesday)$
- relatively complex structures, but clear argument assignments

- Semantic Analysis or Semantic Parsing
- The task of constructing meaning representations from sentences.
- Most popular: Syntax Driven Semantic Analysis

Semantic Analysis

- Semantic Analysis or Semantic Parsing
- The task of constructing meaning representations from sentences.
- Most popular: Syntax Driven Semantic Analysis



[S. G.]

- Based on the principle of compositionality.
 - meaning of the whole built up from the meaning of the parts
 - more specifically, in a way that is guided by word order and syntactic relations.
- Build up the MR by augmenting CFG rules with semantic composition rules.
- Representation produced is literal meaning: context independent and free of inference

- To compute the final MR, we add semantic attachments to our CFG rules.
- These specify how to compute the MR of the parent from those of its children.
- Rules will look like: $A \rightarrow \alpha_1 \dots \alpha_n \{f(\alpha_j.sem, \dots, \alpha_k.sem)\}$
- $A.sem$ (the MR for A) is computed by applying the function f to the MRs of some subset of A 's children.

- AyCaramba serves meat with its parse tree
- Rules with semantic attachments for nouns and NPs:
 - ProperNoun \rightarrow AyCaramba {**AyCaramba**}
 - MassNoun \rightarrow meat {**Meat**}
 - NP \rightarrow ProperNoun {**ProperNoun.sem**}
 - NP \rightarrow MassNoun {**MassNoun.sem**}
- Unary rules normally just copy the semantics of the child to the parents (as in NP rules here).


- we could have: $\lambda y.\lambda x.Serving(x, y)$
- the rule:
Verb \rightarrow serves $\{\lambda y.\lambda x.Serving(x, y)\}$
- or, for **reified events**:
 $\lambda y.\lambda x.\exists e.Serving(e) \wedge Server(e, x) \wedge Served(e, y)$
- the rule:
Verb \rightarrow serves $\{\lambda y.\lambda x.\exists e.Serving(e) \wedge Server(e, x) \wedge Served(e, y)\}$

Building Larger Constituents

- The remaining rules specify how to apply λ -expressions to their arguments.
- For example, the VP rule is:
$$\text{VP} \rightarrow \text{Verb NP} \{ \mathbf{Verb.sem(NP.sem)} \}$$


Building Larger Constituents

- The remaining rules specify how to apply λ -expressions to their arguments.
- For example, the VP rule is:
 $VP \rightarrow \text{Verb NP } \{\mathbf{Verb.sem(NP.sem)}\}$



Building Larger Constituents

- The remaining rules specify how to apply λ -expressions to their arguments.
- For example, the VP rule is:
 $VP \rightarrow \text{Verb NP } \{\mathbf{Verb.sem(NP.sem)}\}$



- $VP.sem = \lambda y. \lambda x. \exists e. Serving(e) \wedge Server(e, x) \wedge Served(e, y)(Meat)$
 $= \lambda x. \exists e. Serving(e) \wedge Server(e, x) \wedge Served(e, Meat)$

Putting Everything Together

- The final rule: $S \rightarrow NP \ VP \ \{\mathbf{VP.sem(NP.sem)}\}$
- with $NP.sem = AyCaramba$

Putting Everything Together

- The final rule: $S \rightarrow NP VP \{ \mathbf{VP.sem(NP.sem)} \}$
- with $NP.sem = AyCaramba$
- then:

$S.sem =$

$$\begin{aligned} & \lambda x. \exists e. Serving(e) \wedge Server(e, x) \wedge Served(e, Meat)(AyCaramba) \\ & = \exists e. Serving(e) \wedge Server(e, AyCaramba) \wedge Served(e, Meat) \end{aligned}$$

Putting Everything Together

- The final rule: $S \rightarrow NP VP \{ \mathbf{VP.sem(NP.sem)} \}$
- with $NP.sem = AyCaramba$
- then:
$$S.sem =$$
$$\lambda x. \exists e. Serving(e) \wedge Server(e, x) \wedge Served(e, Meat)(AyCaramba)$$
$$= \exists e. Serving(e) \wedge Server(e, AyCaramba) \wedge Served(e, Meat)$$
- how about **Every** kid cries
$$\forall x. Child(x) \Rightarrow \exists e. Sleeping(e) \wedge Sleeper(e, x)$$

Putting Everything Together

- The final rule: $S \rightarrow NP VP \{ \mathbf{VP.sem(NP.sem)} \}$
- with $NP.sem = AyCaramba$
- then:
$$S.sem =$$
$$\lambda x. \exists e. Serving(e) \wedge Server(e, x) \wedge Served(e, Meat)(AyCaramba)$$
$$= \exists e. Serving(e) \wedge Server(e, AyCaramba) \wedge Served(e, Meat)$$
- how about **Every** kid cries
$$\forall x. Child(x) \Rightarrow \exists e. Sleeping(e) \wedge Sleeper(e, x)$$
- Is this correct?

- Define S.sem as $\text{NP.sem}(\text{VP.sem})$ as well
- Make NP.sem into a functor by adding λ
 $\lambda Q \forall x. \text{Child}(x) \Rightarrow Q(x)$
- then, apply it to VP.sem:
 $\lambda Q \forall x. \text{Child}(x) \Rightarrow Q(x) (\lambda y. \exists e. \text{Sleeping}(e) \wedge \text{Sleeper}(e, y))$
 $\forall x. \text{Child}(x) \Rightarrow (\lambda y. \exists e. \text{Sleeping}(e) \wedge \text{Sleeper}(e, y))(\mathbf{x})$
 $\forall x. \text{Child}(x) \Rightarrow \exists e. \text{Sleeping}(e) \wedge \text{Sleeper}(e, \mathbf{x})$

The Right NP.sem?

- We will need a new set of noun rules:

Noun \rightarrow Child $\{ \lambda x. Child(x) \}$

Det \rightarrow Every $\{ \lambda P. \lambda Q. \forall x. P(x) \Rightarrow Q(x) \}$

NP \rightarrow Det Noun $\{ \mathbf{Det.sem(Noun.sem)} \}$

- for **Every Child**, we have:

$\lambda P \lambda Q \forall x. P(x) \Rightarrow Q(x) (\lambda x. Child(x))$

$\lambda Q \forall x. (\lambda x. Child(x)) (x) \Rightarrow Q(x)$

$\lambda Q \forall x. Child(x) \Rightarrow Q(x)$

For Proper Nouns

- Assign a different MR to proper nouns, allowing them to take VPs as arguments:

ProperNoun \rightarrow Kate $\{\lambda P. P(\mathbf{Kate})\}$

- For *Kate sleeps*, we have

$\lambda P. P(Kate)(\lambda y. \exists e. Sleeping(e) \wedge Sleeper(e, y))$

$(\lambda y. \exists e. Sleeping(e) \wedge Sleeper(e, y))(\mathbf{Kate})$

$\exists e. Sleeping(e) \wedge Sleeper(e, Kate)$

- Here, we **type-raised** the the argument a of a function f , turning it into a function g that takes f as argument.

- $S \rightarrow NP\ VP\ \{\mathbf{NP.sem(VP.sem)}\}$
- $VP \rightarrow Verb\ \{\mathbf{Verb.sem}\}$
- $VP \rightarrow Verb\ NP\ \{\mathbf{Verb.sem(NP.sem)}\}$
- $NP \rightarrow Det\ Noun\ \{\mathbf{Det.sem(NP.sem)}\}$
- $NP \rightarrow ProperNoun\ \{\mathbf{ProperNoun.sem}\}$
- $Det \rightarrow Every\ \{\lambda P.\lambda Q.\forall x.P(x) \Rightarrow Q(x)\}$
- $Noun \rightarrow Child\ \{\lambda x.Child(x)\}$
- $ProperNoun \rightarrow Kate\ \{\lambda P.P(Kate)\}$
- $Verb \rightarrow sleeps\ (\lambda x.\exists e.Sleeping(e) \wedge Sleeper(e, x))$
- $Verb \rightarrow serves\ \{\lambda y.\lambda x.\exists e.Serving(e) \wedge Server(e, x) \wedge Served(e, y)\}$

- Given a CFG with semantic attachments, how do we obtain the semantic analysis of a sentence?
- One option (integrated): Modify syntactic parser to apply semantic attachments at the time syntactic constituents are constructed.
- Second option (pipelined): Complete the syntactic parse, then walk the tree bottom-up to apply semantic attachments.

- Much current research focuses on learning semantic grammars rather than hand-engineering them.
- Given sentences paired with meaning representations, can we automatically learn
 - Which words are associated with which bits of MR?
 - How those bits combine (in parallel with the syntax) to yield the final MR?
- Second option (pipelined): Complete the syntactic parse, then walk the tree bottom-up to apply semantic attachments.

- Choi, E., Kwiatkowski, T., and Zettlemoyer, L. (2015). Scalable semantic parsing with partial ontologies. In Proceedings of the Association for Computational Linguistics.
- Kwiatkowski, T., Zettlemoyer, L., Goldwater, S., and Steedman, M. (2010). Inducing probabilistic CCG grammars from logical form with higher-order unification. In Proceedings of the Conference on Empirical Methods in Natural Language Processing.
- Reddy, S., Lapata, M., and Steedman, M. (2014). Large-scale semantic parsing without question- answer pairs. Transactions of the Association for Computational Linguistics, 2:377392.
- Zettlemoyer, L. S. and Collins, M. (2005). Learning to map sentences to logical form: Structured classification with probabilistic categorial grammars. In Proceedings of the Conference on Uncertainty in Artificial Intelligence.